

agentize.eu PoC — Teljes Technikai Specifikáció

Verzió: 1.0 **Dátum:** 2026-02-26 **Cél:** Vibe-coding reference — minden információ egy helyen a PoC implementálásához **Időkeret:** 10 munkanap

1. PROJEKT ÁTTEKINTÉS

1.1 Mi ez?

Enterprise AI platform PoC, amely Microsoft Teams-be (és Telegram-ba) integrált AI ügynököt biztosít, az ügyfél Azure előfizetésében futva, VNET izolációval, EU Data Zone Standard telepítéssel.

1.2 PoC Scope — KRISTÁLYTISZTA

BENNE VAN:

- Azure infrastruktúra Bicep template-ból (1 gombnyomás deploy)
- FastAPI backend + LangGraph agent orchestráció
- Teams Bot + Telegram Bot (Bot Framework, multi-channel)
- Adaptive Cards interakció (többpontos jóváhagyás)
- PDF generálás (az elsődleges output formátum)
- Web app PDF letöltéshez (ha SharePoint integráció nem fér bele)
- Cosmos DB state management + LangGraph checkpointing
- Entra ID app registration + SSO
- EU AI Act átláthatósági jelölés minden AI output-on
- Basic VNET izoláció

NINCS BENNE:

- Azure AI Search / RAG — NEM dokumentumokból dolgozunk, hanem user input-ból
- Hallucináció framework (confidence scoring, golden dataset) — a többpontos jóváhagyási workflow kezeli
- Private Link full hardening (7 pontos checklist)
- TISAX dokumentáció
- React Tab szerkesztő (Adaptive Card elég)
- SharePoint mentés (web app-ból letölthető PDF)
- Multi-tenancy
- Metered billing
- Managed Application wrapper (sima resource group + RBAC)

1.3 Kulcsdöntések (Péter feedback alapján)

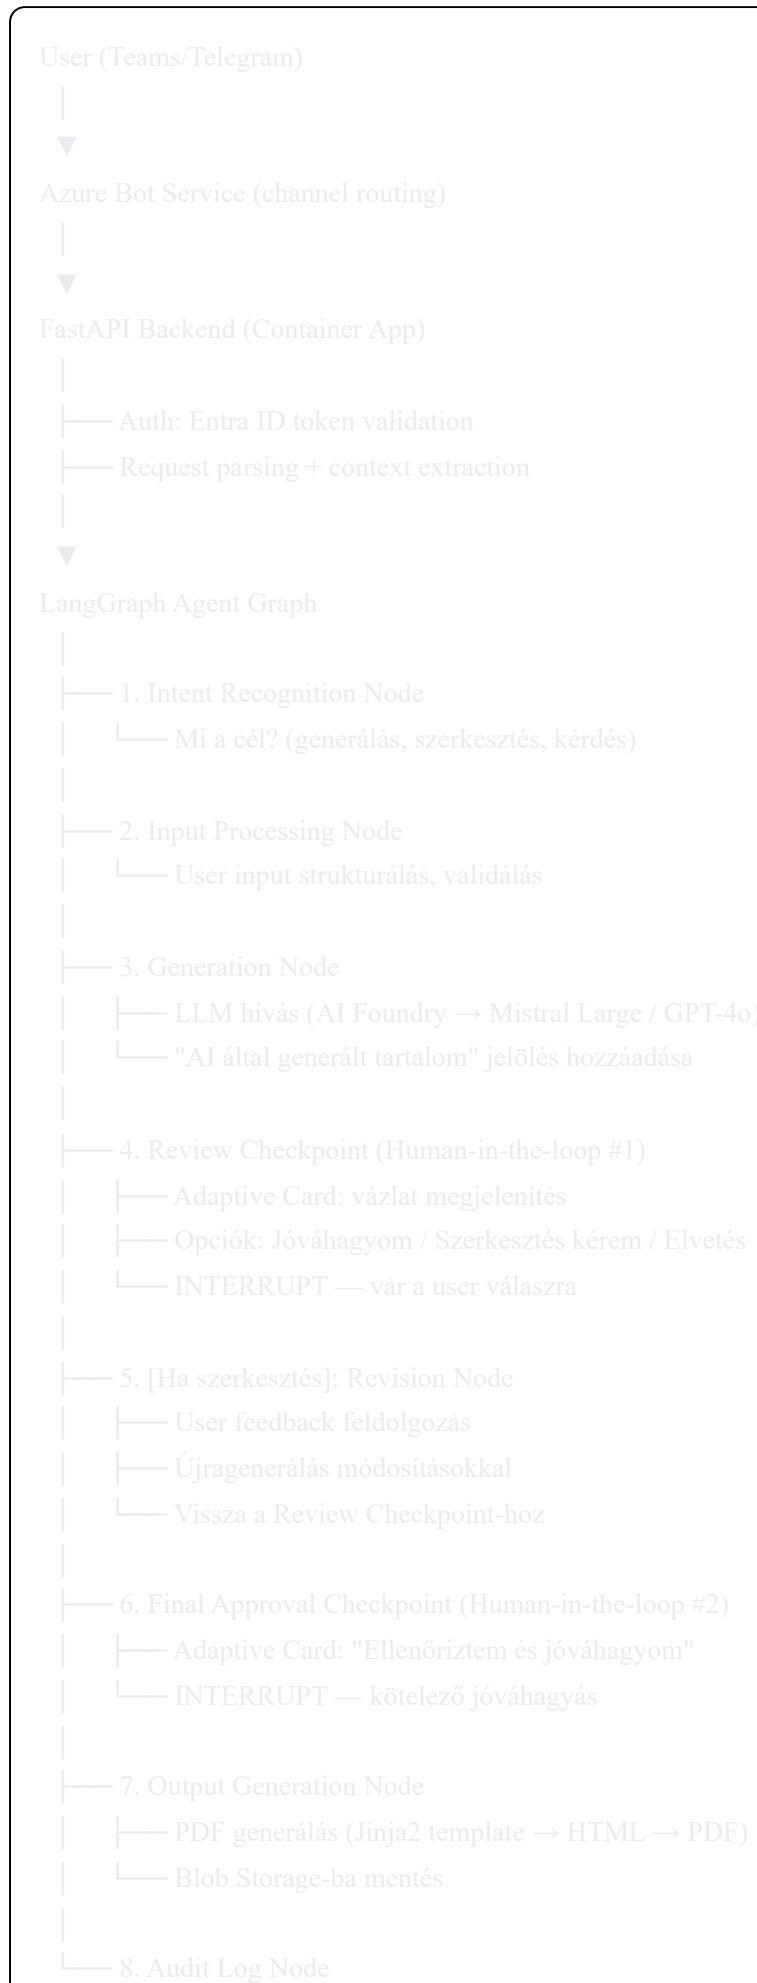
1. **Nincs AI Search** — az agent user input-ból dolgozik, nem dokumentum-RAG-ból. Ha később kell RAG, az "bővített csomag".
2. **Nincs külön hallucináció framework** — folyamatos interakció + többpontos jóváhagyás van beépítve a workflow-ba.
3. **PDF kell** — a jelenlegi outputok PDF-ben készülnek, ez a PoC-ban is kell.
4. **Telegram támogatás** — a Bot Framework natívan kezeli, +1 channel bekapcsolás.
5. **Nem Managed App** — PoC-ban sima resource group + explicit RBAC. IP védelem szerződéses.
6. **Token költség → agent vendor** — a platform infra fix, az LLM fogyasztás az agent vendor dolga.

2. ARCHITEKTÚRA

2.1 High-Level Architecture



2.2 Request Flow



```
|--- Cosmos DB: request, LLM version, approval timestamp, user ID  
└--- Notification: "A dokumentum elkészült: [letöltés link]"
```

2.3 Data Flow — Részletes

1. User → Teams/Telegram: "Készíts egy TWI utasítást a CNC-01 gép beállításáról"

2. Bot Service → FastAPI: POST /api/messages

```
{  
  "type": "message",  
  "text": "Készíts egy TWI utasítást...",  
  "from": { "id": "user-entra-id", "name": "Kovács János" },  
  "channelId": "msteams" | "telegram",  
  "conversation": { "id": "conv-123" }  
}
```

3. FastAPI → LangGraph: invoke graph with state

```
{  
  "user_id": "user-entra-id",  
  "tenant_id": "tenant-123",  
  "channel": "msteams",  
  "message": "Készíts egy TWI utasítást...",  
  "conversation_id": "conv-123",  
  "history": [...previous messages...]  
}
```

4. LangGraph Intent Node → "generate_twi"

5. LangGraph Generation Node → AI Foundry API:

```
POST https://<endpoint>.swedencentral.inference.ai.azure.com/chat/completions  
{  
  "model": "mistral-large-latest",  
  "messages": [  
    {"role": "system", "content": "<TWI system prompt>"},  
    {"role": "user", "content": "Készíts egy TWI utasítást..."}  
  ],  
  "temperature": 0.3,  
  "max_tokens": 4000  
}
```

6. LangGraph Review Checkpoint → INTERRUPT

Bot sends Adaptive Card to user with draft

7. User approves → LangGraph resumes

8. LangGraph Output Node → PDF generation → Blob Storage

Bot sends download link to user

9. Audit log → Cosmos DB

3. INFRASTRUKTÚRA — BICEP SPECIFIKÁCIÓ

3.1 Resource Group Struktúra (PoC)

```
Resource Group: rg-agentize-poc-swedencentral
├── Virtual Network: vnet-agentize-poc
│   ├── Subnet: snet-container-apps (/23, min 256 IPs)
│   ├── Subnet: snet-private-endpoints (/24)
│   └── NSG: nsg-agentize-poc
├── Container App Environment: cae-agentize-poc
│   └── Container App: ca-agentize-backend (FastAPI + LangGraph)
├── Azure AI Foundry: ai-agentize-poc
│   └── Model Deployment: mistral-large (Data Zone Standard)
├── Cosmos DB Account: cosmos-agentize-poc (MongoDB API)
│   └── Database: agentize-poc-db
│       ├── Collection: conversations
│       ├── Collection: agent_state (LangGraph checkpoints)
│       ├── Collection: audit_log
│       └── Collection: generated_documents
├── Storage Account: stagentizepoc
│   └── Container: pdf-output
├── Key Vault: kv-agentize-poc
├── Application Insights: ai-agentize-poc-insights
├── Azure Bot Service: bot-agentize-poc
│   ├── Channel: Microsoft Teams
│   └── Channel: Telegram
└── Entra ID App Registration: app-agentize-poc
```

3.2 Bicep Paraméterek

bicep

```
// main.bicep - Paraméterek
@description('Azure region - KIZÁRÓLAG Sweden Central')
param location string = 'swedencentral'

@description('Projekt prefix')
param projectPrefix string = 'agentize-poc'

@description('AI Foundry model')
@allowed(['mistral-large-latest', 'gpt-4o'])
param aiModel string = 'mistral-large-latest'

@description('Container App min replicas (1 = no cold start)')
param minReplicas int = 1

@description('Cosmos DB throughput mode')
@allowed(['serverless', 'provisioned'])
param cosmosThroughputMode string = 'serverless'

@description('Bot Microsoft App ID (Entra ID-ból)')
param botAppId string

@description('Bot Microsoft App Password (Key Vault-ból)')
@secure()
param botAppPassword string

@description('Telegram Bot Token (opcionális)')
@secure()
param telegramBotToken string = "
```

3.3 VNET Konfiguráció

bicep

```
resource vnet 'Microsoft.Network/virtualNetworks@2023-09-01' = {
  name: 'vnet-${projectPrefix}'
  location: location
  properties: {
    addressSpace: {
      addressPrefixes: ['10.0.0.0/16']
    }
    subnets: [
      {
        name: 'snet-container-apps'
        properties: {
          addressPrefix: '10.0.0.0/23' // Min /23 a Container Apps-hez (256 IP)
          delegations: [
            {
              name: 'Microsoft.App.environments'
              properties: {
                serviceName: 'Microsoft.App/environments'
              }
            }
          ]
        }
      }
    ]
  }
}

{
  name: 'snet-private-endpoints'
  properties: {
    addressPrefix: '10.0.2.0/24'
  }
}
```

3.4 Container App

bicep

```

resource containerAppEnv 'Microsoft.App/managedEnvironments@2023-05-01' = {
  name: 'cae-${projectPrefix}'
  location: location
  properties: {
    vnetConfiguration: {
      infrastructureSubnetId: vnet.properties.subnets[0].id
      internal: true // Csak VNET-en belülről elérhető
    }
    appLogsConfiguration: {
      destination: 'azure-monitor'
    }
  }
}

resource backendApp 'Microsoft.App/containerApps@2023-05-01' = {
  name: 'ca-${projectPrefix}-backend'
  location: location
  properties: {
    managedEnvironmentId: containerAppEnv.id
    configuration: {
      ingress: {
        external: true // Bot Service-nek el kell érnie
        targetPort: 8000
        transport: 'http'
      }
      secrets: [
        { name: 'ai-foundry-key', keyVaultUrl: '${keyVault.properties.vaultUri}secrets/ai-foundry-key' }
        { name: 'cosmos-connection', keyVaultUrl: '${keyVault.properties.vaultUri}secrets/cosmos-connection' }
        { name: 'bot-app-password', keyVaultUrl: '${keyVault.properties.vaultUri}secrets/bot-app-password' }
        { name: 'blob-connection', keyVaultUrl: '${keyVault.properties.vaultUri}secrets/blob-connection' }
      ]
    }
    template: {
      containers: [
        {
          name: 'backend'
          image: 'gher.io/agentize-eu/poc-backend:latest'
          resources: {
            cpu: json('1.0')
            memory: '2Gi'
          }
          env: [
            { name: 'AI_FOUNDRY_ENDPOINT', value: '<AI Foundry endpoint URL>' }
            { name: 'AI_FOUNDRY_KEY', secretRef: 'ai-foundry-key' }
            { name: 'COSMOS_CONNECTION', secretRef: 'cosmos-connection' }
            { name: 'BOT_APP_ID', value: botAppId }
          ]
        }
      ]
    }
  }
}

```

```
{ name: 'BOT_APP_PASSWORD', secretRef: 'bot-app-password' }
{ name: 'BLOB_CONNECTION', secretRef: 'blob-connection' }
{ name: 'APPLICATIONINSIGHTS_CONNECTION_STRING', value: appInsights.properties.ConnectionString }

]
}

]

scale: {
  minReplicas: minReplicas // 1 = nincs cold start (~$10/hó)
  maxReplicas: 5
  rules: [
    {
      name: 'http-scaling'
      http: { metadata: { concurrentRequests: '20' } }
    }
  ]
}

}

}

}

}

}
```

3.5 Azure AI Foundry

bicep

```

// Azure AI Foundry (korábban AI Model Catalog)
// KRITIKUS: Data Zone Standard, Sweden Central
resource aiFoundry 'Microsoft.CognitiveServices/accounts@2024-04-01-preview' = {
    name: 'ai-${projectPrefix}'
    location: location // swedencentral
    kind: 'AIServices'
    sku: { name: 'S0' }
    properties: {
        customSubDomainName: 'ai-${projectPrefix}'
        publicNetworkAccess: 'Enabled' // PoC-ban OK, MVP-ben Disabled + Private Endpoint
        // Data Zone Standard — EU adatrezidencia garancia
        // FONTOS: Csak Sweden Central és Germany West Central ad szerződéses EU garanciát
    }
}

resource aiDeployment 'Microsoft.CognitiveServices/accounts/deployments@2024-04-01-preview' = {
    parent: aiFoundry
    name: aiModel
    sku: {
        name: 'DataZoneStandard' // NEM Global Standard!
        capacity: 10 // TPM (tokens per minute) in thousands
    }
    properties: {
        model: {
            format: 'MistralAI' // vagy 'OpenAI' ha GPT-4o
            name: aiModel
            version: 'latest'
        }
    }
}

```

3.6 Cosmos DB

bicep

```

resource cosmosAccount 'Microsoft.DocumentDB/databaseAccounts@2023-11-15' = {
  name: 'cosmos-${projectPrefix}'
  location: location
  kind: 'MongoDB'
  properties: {
    databaseAccountOfferType: 'Standard'
    locations: [{ locationName: location, failoverPriority: 0 }]
    capabilities: cosmosThroughputMode == 'serverless'
      ? [{ name: 'EnableServerless' }, { name: 'EnableMongo' }]
      : [{ name: 'EnableMongo' }]
    publicNetworkAccess: 'Enabled' // PoC-ban OK
    // MongoDB API — kompatibilis a meglévő kódbázissal
  }
}

resource cosmosDb 'Microsoft.DocumentDB/databaseAccounts/mongodbDatabases@2023-11-15' = {
  parent: cosmosAccount
  name: '${projectPrefix}-db'
  properties: {
    resource: { id: '${projectPrefix}-db' }
  }
}

// Collections — lásd a 6. szekciót a sémához

```

3.7 Blob Storage + Key Vault + App Insights

bicep

```

resource storageAccount 'Microsoft.Storage/storageAccounts@2023-01-01' = {
  name: replace('st${projectPrefix}', '_', '')
  location: location
  sku: { name: 'Standard_LRS' }
  kind: 'StorageV2'
  properties: {
    publicNetworkAccess: 'Enabled' // PoC-ban OK
    allowBlobPublicAccess: false
    minimumTlsVersion: 'TLS1_2'
  }
}

resource blobContainer 'Microsoft.Storage/storageAccounts/blobServices/containers@2023-01-01' = {
  name: '${storageAccount.name}/default/pdf-output'
  properties: { publicAccess: 'None' }
}

resource keyVault 'Microsoft.KeyVault/vaults@2023-07-01' = {
  name: 'kv-${projectPrefix}'
  location: location
  properties: {
    sku: { family: 'A', name: 'standard' }
    tenantId: subscription().tenantId
    enableRbacAuthorization: true
    publicNetworkAccess: 'Enabled' // PoC-ban OK
  }
}

resource appInsights 'Microsoft.Insights/components@2020-02-02' = {
  name: 'ai-${projectPrefix}-insights'
  location: location
  kind: 'web'
  properties: {
    Application_Type: 'web'
    RetentionInDays: 30
  }
}

```

3.8 Azure Bot Service

bicep

```

resource botService 'Microsoft.BotService/botServices@2022-09-15' = {
    name: 'bot-${projectPrefix}'
    location: 'global' // Bot Service minden global
    kind: 'azurebot'
    sku: { name: 'S1' }
    properties: {
        displayName: 'agentize.eu PoC Bot'
        description: 'Enterprise AI Platform PoC'
        endpoint: 'https://${backendApp.properties.configuration.ingress.fqdn}/api/messages'
        msaAppId: botAppId
        msaAppType: 'SingleTenant'
        msaAppTenantId: subscription().tenantId
    }
}

// Teams channel
resource teamsChannel 'Microsoft.BotService/botServices/channels@2022-09-15' = {
    parent: botService
    name: 'MsTeamsChannel'
    location: 'global'
    properties: {
        channelName: 'MsTeamsChannel'
        properties: { isEnabled: true }
    }
}

// Telegram channel (opcionális)
resource telegramChannel 'Microsoft.BotService/botServices/channels@2022-09-15' = if (telegramBotToken != "") {
    parent: botService
    name: 'TelegramChannel'
    location: 'global'
    properties: {
        channelName: 'TelegramChannel'
        properties: {
            accessToken: telegramBotToken
            isEnabled: true
        }
    }
}

```

4. BACKEND — FASTAPI + LANGGRAPH

4.1 Projekt Struktúra

```
poc-backend/
├── Dockerfile
├── requirements.txt
├── pyproject.toml
├── .env.example
├── .devcontainer/
│   └── devcontainer.json
|
└── app/
    ├── __init__.py
    ├── main.py          # FastAPI app + Bot endpoint
    ├── config.py        # Environment config (pydantic-settings)
    |
    ├── bot/
    │   ├── __init__.py
    │   ├── bot_handler.py      # Bot Framework message handler
    │   ├── adaptive_cards.py  # Adaptive Card JSON templates
    │   └── teams_helpers.py   # Teams/Telegram specifikus logika
    |
    ├── agent/
    │   ├── __init__.py
    │   ├── graph.py         # LangGraph graph definition (A FŐ LOGIKA)
    │   ├── nodes/
    │   │   ├── __init__.py
    │   │   ├── intent.py      # Intent recognition node
    │   │   ├── process_input.py # User input processing node
    │   │   ├── generate.py    # LLM generation node
    │   │   ├── review.py      # Human-in-the-loop review checkpoint
    │   │   ├── revise.py      # Revision node (user feedback)
    │   │   ├── approve.py     # Final approval checkpoint
    │   │   ├── output.py      # PDF generation + storage
    │   │   └── audit.py       # Audit logging node
    │   ├── state.py         # LangGraph State definition
    │   ├── prompts/
    │   │   ├── twi_system.txt # TWI system prompt
    │   │   ├── twi_generate.txt # TWI generation prompt template
    │   │   └── intent_classify.txt # Intent classification prompt
    │   └── tools/
    │       ├── __init__.py
    │       └── pdf_generator.py # PDF generation tool
    |
    └── services/
```

```
| | | └── __init__.py
| | |   └── ai_foundry.py      # Azure AI Foundry client
| | |   └── cosmos_db.py      # Cosmos DB MongoDB client
| | |   └── blob_storage.py   # Blob Storage client (PDF upload)
| | |   └── key_vault.py      # Key Vault secret client
|
| | |
| | └── models/
| |   └── __init__.py
| |   └── conversation.py    # Conversation data model
| |   └── twi_document.py     # TWI document model
| |   └── audit_entry.py     # Audit log entry model
|
| |
| └── templates/
|   └── twi_template.html    # Jinja2 HTML template → PDF
|   └── twi_style.css        # PDF styling
|
| └── tests/
|   └── test_graph.py
|   └── test_generation.py
|   └── test_pdf.py
|
└── infra/
  └── main.bicep
  └── parameters.json
  └── deploy.sh
```

4.2 Requirements

txt

```
# requirements.txt

# Web framework
fastapi==0.115.*
uvicorn[standard]==0.32.*

# Bot Framework
botbuilder-core==4.16.*
botbuilder-integration-aiohttp==4.16.*
aiohttp==3.10.*

# LangGraph + LangChain (CSAK LangGraph-ot használunk)
langgraph==0.3.*
langchain-core==0.3.*
langsmith==0.2.*

# Azure services
azure-identity==1.19.*
azure-keyvault-secrets==4.9.*
azure-storage-blob==12.23.*
azure-ai-inference==1.0.*      # AI Foundry client
pymongo==4.10.*              # Cosmos DB MongoDB API
motor==3.6.*                  # Async MongoDB driver

# PDF generation
jinja2==3.1.*
weasyprint==63.*             # HTML → PDF

# Config & utils
pydantic==2.10.*
pydantic-settings==2.7.*
python-dotenv==1.0.*

# Observability
opentelemetry-api==1.29.*
opentelemetry-sdk==1.29.*
azure-monitor-opentelemetry==1.6.*

# Dev/test
pytest==8.3.*
pytest-asyncio==0.25.*
httpx==0.28.*                 # Async HTTP test client
```

4.3 Config

python

```
# app/config.py
from pydantic_settings import BaseSettings

class Settings(BaseSettings):
    # Azure AI Foundry
    ai_foundry_endpoint: str
    ai_foundry_key: str
    ai_model: str = "mistral-large-latest"
    ai_temperature: float = 0.3
    ai_max_tokens: int = 4000

    # Cosmos DB
    cosmos_connection: str
    cosmos_database: str = "agentize-poc-db"

    # Blob Storage
    blob_connection: str
    blob_container: str = "pdf-output"

    # Bot Framework
    bot_app_id: str
    bot_app_password: str

    # Application Insights
    applicationinsights_connection_string: str = ""

    # App config
    environment: str = "poc"
    log_level: str = "INFO"

    class Config:
        env_file = ".env"
        case_sensitive = False

settings = Settings()
```

4.4 FastAPI Main

```
python
```

```
# app/main.py
import logging
from fastapi import FastAPI, Request, Response
from botbuilder.core import BotFrameworkAdapter, BotFrameworkAdapterSettings
from botbuilder.schema import Activity

from app.config import settings
from app.bot.bot_handler import AgentizeBotHandler

# Logging
logging.basicConfig(level=settings.log_level)
logger = logging.getLogger(__name__)

# FastAPI app
app = FastAPI(
    title="agentize.eu PoC Backend",
    version="0.1.0",
    docs_url="/docs" if settings.environment == "poc" else None,
)

# Bot Framework adapter
adapter_settings = BotFrameworkAdapterSettings(
    app_id=settings.bot_app_id,
    app_password=settings.bot_app_password,
)
adapter = BotFrameworkAdapter(adapter_settings)

# Bot handler
bot = AgentizeBotHandler()

# Error handler
async def on_error(context, error):
    logger.error(f"Bot error: {error}", exc_info=True)
    await context.send_activity("Hiba történt. Kérlek próbáld újra.")

adapter.on_turn_error = on_error

@app.post("/api/messages")
async def messages(request: Request):
    """Bot Framework messaging endpoint"""
    body = await request.json()
    activity = Activity().deserialize(body)
    auth_header = request.headers.get("Authorization", "")

    response = await adapter.process_activity(
```

```
    activity, auth_header, bot.on_turn
)

if response:
    return Response(
        content=response.body,
        status_code=response.status,
        headers=dict(response.headers) if response.headers else {},
    )
return Response(status_code=200)
```

```
@app.get("/health")
async def health():
    return {"status": "healthy", "environment": settings.environment}
```

```
@app.get("/")
async def root():
    return {"service": "agentize.eu PoC Backend", "version": "0.1.0"}
```

4.5 Bot Handler

```
python
```

```
# app/bot/bot_handler.py
import logging
from botbuilder.core import ActivityHandler, TurnContext, CardFactory
from botbuilder.schema import ActivityTypes, Attachment

from app.agent.graph import create_agent_graph, run_agent
from app.bot.adaptive_cards import (
    create_review_card,
    create_approval_card,
    create_result_card,
    create_welcome_card,
)
from app.services.cosmos_db import ConversationStore

logger = logging.getLogger(__name__)

class AgentizeBotHandler(ActivityHandler):
    def __init__(self):
        self.graph = create_agent_graph()
        self.conversation_store = ConversationStore()

    async def on_message_activity(self, turn_context: TurnContext):
        """Bejövő üzenet feldolgozás"""
        user_id = turn_context.activity.from_property.id
        conversation_id = turn_context.activity.conversation.id
        channel_id = turn_context.activity.channel_id
        text = turn_context.activity.text or ""
        value = turn_context.activity.value # Adaptive Card submit value

        logger.info(f"Message from {user_id} on {channel_id}: {text[:50]}...")

        # Adaptive Card válasz (jóváhagyás / szerkesztés kérés / elvetés)
        if value:
            await self._handle_card_action(turn_context, value, conversation_id, user_id)
            return

        # Normál szöveges üzenet → Agent graph futtatás
        await self._handle_text_message(turn_context, text, conversation_id, user_id, channel_id)

    async def _handle_text_message(self, turn_context, text, conversation_id, user_id, channel_id):
        """Szöveges üzenet → LangGraph agent"""
        # "Gondolkozom..." jelzés
        await turn_context.send_activity("⏳ Feldolgozom a kéréseket...")

        # Agent futtatás
```

```
result = await run_agent(  
    graph=self.graph,  
    message=text,  
    user_id=user_id,  
    conversation_id=conversation_id,  
    channel=channel_id,  
)  
  
# Eredmény kezelése  
if result["status"] == "review_needed":  
    # Vázlat kész → review Adaptive Card küldése  
    card = create_review_card(  
        draft=result["draft"],  
        metadata=result["metadata"],  
    )  
    await turn_context.send_activity(  
        Activity(  
            type=ActivityTypes.message,  
            attachments=[CardFactory.adaptive_card(card)],  
        )  
    )  
  
elif result["status"] == "clarification_needed":  
    # Nem egyértelmű input → kérdés vissza  
    await turn_context.send_activity(result["message"])  
  
elif result["status"] == "error":  
    await turn_context.send_activity(f"❌ Hiba: {result['message']}")  
  
async def _handle_card_action(self, turn_context, value, conversation_id, user_id):  
    """Adaptive Card válasz feldolgozás"""  
    action = value.get("action")  
  
    if action == "approve_draft":  
        # Review #1 — vázlat jóváhagyva → Final approval card  
        card = create_approval_card(  
            draft=value.get("draft"),  
            metadata=value.get("metadata"),  
        )  
        await turn_context.send_activity(  
            Activity(  
                type=ActivityTypes.message,  
                attachments=[CardFactory.adaptive_card(card)],  
            )  
        )  
  
    elif action == "request_edit":
```

```

# Szerkesztés kérés → Revision node
feedback = value.get("feedback", "")
await turn_context.send_activity("⏳ Módosítom a szerkesztési kérésed alapján...")

result = await run_agent(
    graph=self.graph,
    message=feedback,
    user_id=user_id,
    conversation_id=conversation_id,
    resume_from="revision",
    context=value,
)

card = create_review_card(
    draft=result["draft"],
    metadata=result["metadata"],
)
await turn_context.send_activity(
    Activity(
        type=ActivityTypes.message,
        attachments=[CardFactory.adaptive_card(card)],
    )
)

elif action == "final_approve":
    # Final approval → PDF generálás
    await turn_context.send_activity("⏳ PDF generálás folyamatban...")

result = await run_agent(
    graph=self.graph,
    message="",
    user_id=user_id,
    conversation_id=conversation_id,
    resume_from="output",
    context=value,
)

card = create_result_card(
    pdf_url=result["pdf_url"],
    document_title=result["title"],
    metadata=result["metadata"],
)
await turn_context.send_activity(
    Activity(
        type=ActivityTypes.message,
        attachments=[CardFactory.adaptive_card(card)],
    )
)

```

```
)  
  
elif action == "reject":  
    await turn_context.send_activity("⚠️ Elvettem a vázlatot. Új kéréssel indíthatsz újat.")  
  
async def on_members_added_activity(self, members_added, turn_context: TurnContext):  
    """Üdvözlő üzenet új member-nek"""  
    for member in members_added:  
        if member.id != turn_context.activity.recipient.id:  
            card = create_welcome_card()  
            await turn_context.send_activity(  
                Activity(  
                    type=ActivityTypes.message,  
                    attachments=[CardFactory.adaptive_card(card)],  
                )  
            )
```

4.6 LangGraph Agent Graph

python

```
# app/agent/state.py
from typing import TypedDict, Literal, Optional, List, Any
from langgraph.graph import MessagesState

class AgentState(TypedDict):
    """LangGraph agent state"""
    # Input
    user_id: str
    tenant_id: str
    conversation_id: str
    channel: str
    message: str

    # Processing
    intent: Optional[str] # "generate_twi", "edit_twi", "question", "unknown"
    processed_input: Optional[dict]
    draft: Optional[str]
    draft_metadata: Optional[dict]

    # Revision loop
    revision_feedback: Optional[str]
    revision_count: int

    # Output
    status: str # "review_needed", "approved", "completed", "error"
    pdf_url: Optional[str]
    pdf_blob_name: Optional[str]

    # Audit
    llm_model: Optional[str]
    llm_tokens_used: Optional[int]
    approval_timestamp: Optional[str]

    # History
    messages: List[Any]
```

python

```
# app/agent/graph.py
import logging
from langgraph.graph import StateGraph, END
from langgraph.checkpoint.memory import MemorySaver # PoC-ban in-memory, prod-ban Cosmos DB

from app.agent.state import AgentState
from app.agent.nodes.intent import intent_node
from app.agent.nodes.process_input import process_input_node
from app.agent.nodes.generate import generate_node
from app.agent.nodes.review import review_node
from app.agent.nodes.revise import revise_node
from app.agent.nodes.approve import approve_node
from app.agent.nodes.output import output_node
from app.agent.nodes.audit import audit_node

logger = logging.getLogger(__name__)

def should_generate(state: AgentState) -> str:
    """Intent alapján routing"""
    intent = state.get("intent", "unknown")
    if intent in ("generate_twi", "edit_twi"):
        return "process_input"
    elif intent == "question":
        return "generate" # Egyszerű Q&A, nincs input processing
    else:
        return "clarify" # Nem egyértelmű → kérdezzünk vissza

def after_review(state: AgentState) -> str:
    """Review után routing"""
    status = state.get("status", "")
    if status == "approved":
        return "approve"
    elif status == "revision_requested":
        return "revise"
    else:
        return "end" # rejected

def after_revision(state: AgentState) -> str:
    """Revision után routing — max 3 kör"""
    if state.get("revision_count", 0) >= 3:
        return "approve" # Forced approval after 3 rounds
    return "review"
```

```

def create_agent_graph() -> StateGraph:
    """LangGraph agent gráf felépítés"""

    # Gráf definíció
    builder = StateGraph(AgentState)

    # Node-ok hozzáadása
    builder.add_node("intent", intent_node)
    builder.add_node("process_input", process_input_node)
    builder.add_node("generate", generate_node)
    builder.add_node("review", review_node)      # Human-in-the-loop #1
    builder.add_node("revise", revise_node)
    builder.add_node("approve", approve_node)     # Human-in-the-loop #2
    builder.add_node("output", output_node)       # PDF generation
    builder.add_node("audit", audit_node)

    # Edge-ek (flow)
    builder.set_entry_point("intent")
    builder.add_conditional_edges("intent", should_generate)
    builder.add_edge("process_input", "generate")
    builder.add_edge("generate", "review")

    # Review → conditional
    builder.add_conditional_edges("review", after_review)

    # Revision loop
    builder.add_edge("revise", "generate") # Újragenerálás a módosításokkal

    # Approve → Output → Audit → END
    builder.add_edge("approve", "output")
    builder.add_edge("output", "audit")
    builder.add_edge("audit", END)

    # Checkpointer (PoC: in-memory; Prod: Cosmos DB)
    checkpointer = MemorySaver()

    return builder.compile(
        checkpointer=checkpointer,
        interrupt_before=["review", "approve"], # Human-in-the-loop breakpoints
    )

async def run_agent(
    graph,
    message: str,
    user_id: str,

```

```
conversation_id: str,  
channel: str = "msteams",  
resume_from: str = None,  
context: dict = None,  
) -> dict:  
    """Agent futtatás vagy folytatás"""  
  
    config = {"configurable": {"thread_id": conversation_id}}  
  
    if resume_from:  
        # Folytatás interrupt utánról  
        state_update = _build_resume_state(resume_from, context)  
        result = await graph.invoke(state_update, config)  
    else:  
        # Új futtatás  
        initial_state = AgentState(  
            user_id=user_id,  
            tenant_id="poc-tenant", # PoC-ban fix  
            conversation_id=conversation_id,  
            channel=channel,  
            message=message,  
            revision_count=0,  
            status="processing",  
            messages=[],  
)  
        result = await graph.invoke(initial_state, config)  
  
    return result
```

```
def _build_resume_state(resume_from: str, context: dict) -> dict:  
    """Interrupt utáni state update"""  
    if resume_from == "revision":  
        return {  
            "status": "revision_requested",  
            "revision_feedback": context.get("feedback", "")  
        }  
    elif resume_from == "output":  
        return {  
            "status": "approved",  
            "approval_timestamp": context.get("timestamp")  
        }  
    return {}
```

4.7 Agent Node-ok — Implementáció

python

```
# app/agent/nodes/intent.py
from app.agent.state import AgentState
from app.services.ai_foundry import call_llm
```

INTENT_PROMPT = """Te az agentize.eu AI platform intent felismerő modulja vagy.

Osztályozd a felhasználó kérését az alábbi kategóriák egyikébe:

- generate_twi: Új TWI (Training Within Industry) utasítás generálása
- edit_twi: Meglévő TWI szerkesztése, módosítása
- question: Általános kérdés a rendszerről vagy a folyamatokról
- unknown: Nem egyértelmű, kérdezzünk vissza

VÁLASZOLJ KIZÁRÓLAG az intent nevével, semmi mással.

Felhasználó üzenete: {message}"""

```
async def intent_node(state: AgentState) -> AgentState:
    """Szándék felismerés"""
    response = await call_llm(
        prompt=INTENT_PROMPT.format(message=state["message"]),
        temperature=0.1,
        max_tokens=20,
    )
    intent = response.strip().lower()
    if intent not in ("generate_twi", "edit_twi", "question", "unknown"):
        intent = "unknown"

    return {"**state, "intent": intent}
```

python

```
# app/agent/nodes/generate.py
from app.agent.state import AgentState
from app.services.ai_foundry import call_llm
```

TWI_SYSTEM_PROMPT = """Te az agentize.eu TWI (Training Within Industry) generátor modulja vagy.

FELADATOD:

A felhasználó inputja alapján strukturált munkautasítást generál sz az alábbi formátumban:

1. CÍM: A munkautasítás rövid címe
2. CÉL: Mit ér el a dolgozó, ha követi az utasítást
3. SZÜKSÉGES ANYAGOK ÉS ESZKÖZÖK: Felsorolás
4. BIZTONSÁGI ELŐÍRÁSOK: Releváns biztonsági figyelmeztetések
5. LÉPÉSEK: Számoszott lépések, mindegyikhez:
 - Fölépés: Mit csinálj
 - KulcsPontok: Hogyan csinálj (részletek, amik a minőséget biztosítják)
 - Indoklás: Miért fontos ez a lépés
6. MINŐSÉGI ELLENŐRZÉS: Hogyan ellenőrizhető a munka minősége

SZABÁLYOK:

- minden output AUTOMATIKUSAN tartalmazza: "⚠️ AI által generált tartalom — emberi felülvizsgálat szükséges."
- Légy precíz és konkrét — gyártási környezetben használják
- Ha nem kapsz elég információt, KÉRDEZZ VISSZA — ne találj ki részleteket
- Magyar nyelven válaszolj, technikai szakkifejezések angolul is megadhatók zárójelben

"""

TWI_GENERATE_PROMPT = """A felhasználó kérése:

{message}

{revision_context}

Generáld a TWI utasítást a megadott formátumban."""

```
async def generate_node(state: AgentState) -> AgentState:
```

"""TWI generálás LLM-mel"""
 revision_context = """
 if state.get("revision_feedback"):
 revision_context += f"""
 KORÁBBI VÁZLAT:
 {state.get('draft', '')}

FELHASZNÁLÓI VISSZAJELZÉS:

{state['revision_feedback']}

Módosítsd a vázlatot a visszajelzés alapján.

"""

```
response = await call_llm(  
    system_prompt=TWI_SYSTEM_PROMPT,  
    prompt=TWI_GENERATE_PROMPT.format(  
        message=state["message"],  
        revision_context=revision_context,  
    ),  
    temperature=0.3,  
    max_tokens=4000,  
)  
  
# AI jelölés hozzáadása  
draft = f"⚠️ AI által generált tartalom — emberi felülvizsgálat szükséges.\n\n{response}"  
  
return {  
    **state,  
    "draft": draft,  
    "draft_metadata": {  
        "model": state.get("llm_model", "mistral-large"),  
        "generated_at": _now_iso(),  
        "revision": state.get("revision_count", 0),  
    },  
    "status": "review_needed",  
}
```

python

```
# app/agent/nodes/output.py
import uuid
from app.agent.state import AgentState
from app.agent.tools.pdf_generator import generate_twi_pdf
from app.services.blob_storage import upload_pdf

async def output_node(state: AgentState) -> AgentState:
    """PDF generálás és feltöltés Blob Storage-ra"""
    # PDF generálás
    pdf_bytes = await generate_twi_pdf(
        content=state["draft"],
        metadata=state["draft_metadata"],
        user_id=state["user_id"],
    )

    # Blob Storage feltöltés
    blob_name = f"twi/{state['conversation_id']}/{uuid.uuid4().hex}.pdf"
    pdf_url = await upload_pdf(pdf_bytes, blob_name)

    return {
        **state,
        "pdf_url": pdf_url,
        "pdf_blob_name": blob_name,
        "status": "completed",
    }
```

python

```
# app/agent/nodes/audit.py
from datetime import datetime, timezone
from app.agent.state import AgentState
from app.services.cosmos_db import AuditStore

async def audit_node(state: AgentState) -> AgentState:
    """Audit log mentés Cosmos DB-be"""
    audit_store = AuditStore()
    await audit_store.log({
        "conversation_id": state["conversation_id"],
        "user_id": state["user_id"],
        "tenant_id": state["tenant_id"],
        "channel": state["channel"],
        "intent": state["intent"],
        "llm_model": state.get("llm_model"),
        "revision_count": state.get("revision_count", 0),
        "pdf_blob_name": state.get("pdf_blob_name"),
        "status": state["status"],
        "approval_timestamp": state.get("approval_timestamp"),
        "created_at": datetime.now(timezone.utc).isoformat(),
    })
    return state
```

5. ADAPTIVE CARDS

5.1 Review Card (Human-in-the-loop #1)

```
python
```

```
# app/bot/adaptive_cards.py

def create_review_card(draft: str, metadata: dict) -> dict:
    """Vázlat review card — jóváhagyás / szerkesztés kérés / elvetés"""
    return {
        "$schema": "http://adaptivecards.io/schemas/adaptive-card.json",
        "type": "AdaptiveCard",
        "version": "1.4",
        "body": [
            {
                "type": "TextBlock",
                "text": "📋 TWI Vázlat — Felülvizsgálat szükséges",
                "weight": "bolder",
                "size": "large",
                "wrap": True,
            },
            {
                "type": "TextBlock",
                "text": f"⚠️ AI által generált tartalom | Modell: {metadata.get('model', 'N/A')} | "
                f"Generálva: {metadata.get('generated_at', 'N/A')}",
                "size": "small",
                "color": "warning",
                "wrap": True,
            },
            {"type": "TextBlock", "text": "---", "separator": True},
            {
                "type": "TextBlock",
                "text": draft[:2000], # Adaptive Card limit
                "wrap": True,
                "fontType": "default",
            },
            {
                "type": "TextBlock",
                "text": "---",
                "separator": True,
            },
            {
                "type": "TextBlock",
                "text": "Szerkesztési megjegyzés (opcionális):",
                "size": "small",
            },
            {
                "type": "Input.Text",
                "id": "feedback",
                "isMultiline": True,
                "placeholder": "Pl.: A 3. lépésben hiányzik a hőmérséklet beállítás...",
            }
        ]
    }
```

```
        },
    ],
    "actions": [
        {
            "type": "Action.Submit",
            "title": "✓ Jóváhagyom a vázlatot",
            "style": "positive",
            "data": {
                "action": "approve_draft",
                "draft": draft,
                "metadata": metadata,
            },
        },
        {
            "type": "Action.Submit",
            "title": "📝 Szerkesztés kérem",
            "data": {
                "action": "request_edit",
                "draft": draft,
                "metadata": metadata,
            },
        },
        {
            "type": "Action.Submit",
            "title": "☒ Elvetés",
            "style": "destructive",
            "data": {"action": "reject"},
        },
    ],
}
```

```
def create_approval_card(draft: str, metadata: dict) -> dict:
    """Végső jóváhagyás card (Human-in-the-loop #2) — KÖTELEZŐ"""
    return {
        "$schema": "http://adaptivecards.io/schemas/adaptive-card.json",
        "type": "AdaptiveCard",
        "version": "1.4",
        "body": [
            {
                "type": "TextBlock",
                "text": "🔒 Vélegesítés — Kötelező Jóváhagyás",
                "weight": "bolder",
                "size": "large",
                "color": "attention",
                "wrap": True,
            },
        ],
    }
```

```

{
    "type": "TextBlock",
    "text": "⚠ Ez a dokumentum AI által generált tartalom. "
        "Kérlek ellenőrizd a tartalmat, mielőtt véglegesíted. "
        "Véglegesítés után PDF készül és archiválásra kerül.",
    "wrap": True,
    "color": "warning",
},
{
    "type": "TextBlock",
    "text": draft[:2000],
    "wrap": True,
},
],
"actions": [
{
    "type": "Action.Submit",
    "title": "✓ Ellenőrztem és jóváhagyom",
    "style": "positive",
    "data": {
        "action": "final_approve",
        "draft": draft,
        "metadata": metadata,
        "timestamp": "__CURRENT_TIMESTAMP__"
    },
},
{
    "type": "Action.Submit",
    "title": "➡ Vissza a szerkesztéshez",
    "data": {
        "action": "request_edit",
        "draft": draft,
        "metadata": metadata
    },
},
],
}

```

```

def create_result_card(pdf_url: str, document_title: str, metadata: dict) -> dict:
    """Eredmény card — PDF letöltés link"""
    return {
        "$schema": "http://adaptivecards.io/schemas/adaptive-card.json",
        "type": "AdaptiveCard",
        "version": "1.4",
        "body": [
{

```

```
"type": "TextBlock",
"text": "☑ Dokumentum elkészült",
"weight": "bolder",
"size": "large",
"color": "good",
},
{
"type": "FactSet",
"facts": [
{"title": "Cím:", "value": document_title},
 {"title": "Formátum:", "value": "PDF"},
 {"title": "Modell:", "value": metadata.get("model", "N/A")},
 {"title": "Jóváhagyta:", "value": metadata.get("approved_by", "N/A")},
],
},
],
"actions": [
{
"type": "Action.OpenUrl",
"title": "📥 PDF letöltés",
"url": pdf_url,
},
],
}
```

```
def create_welcome_card() -> dict:
    """Üdvözlő card"""
    return {
        "$schema": "http://adaptivecards.io/schemas/adaptive-card.json",
        "type": "AdaptiveCard",
        "version": "1.4",
        "body": [
            {
                "type": "TextBlock",
                "text": "👋 Üdvözöllek! Én az agentize.eu AI asszisztens vagyok.",
                "weight": "bolder",
                "size": "medium",
                "wrap": True,
            },
            {
                "type": "TextBlock",
                "text": "Segíthetek TWI (Training Within Industry) munkautasítások "
                    "generálásában. Írd le, milyen utasításra van szükséged!",
                "wrap": True,
            },
        ],
    }
```

```
        "type": "TextBlock",
        "text": "Példa: \"Készíts egy TWI utasítást a CNC-01 gép beállításáról\"",
        "wrap": True,
        "isSubtle": True,
        "fontType": "monospace",
    },
],
}
```

6. ADATBÁZIS SÉMA — COSMOS DB (MongoDB API)

6.1 Collections

```
javascript
```

```
// Collection: conversations
{
  "_id": ObjectId,
  "conversation_id": "conv-123",           // Bot Framework conversation ID
  "user_id": "user-entra-id",
  "tenant_id": "poc-tenant",
  "channel": "msteams" | "telegram",
  "started_at": ISODate,
  "last_activity": ISODate,
  "message_count": 5,
  "status": "active" | "completed" | "expired"
}
// Index: { conversation_id: 1 } unique
// Index: { tenant_id: 1, user_id: 1 }
// TTL index: { last_activity: 1 }, expireAfterSeconds: 7776000 (90 nap)
```

```
// Collection: agent_state (LangGraph checkpoints)
// AUTOMATIKUSAN KEZELI a LangGraph checkpointer
// Partition key: thread_id (= conversation_id)
{
  "_id": ObjectId,
  "thread_id": "conv-123",
  "checkpoint_id": "cp-456",
  "parent_checkpoint_id": "cp-455" | null,
  "checkpoint": { /* LangGraph internal state */ },
  "metadata": {
    "step": 3,
    "node": "review",
    "writes": {...}
  },
  "created_at": ISODate
}
```

```
// Collection: generated_documents
{
  "_id": ObjectId,
  "document_id": "doc-uuid",
  "conversation_id": "conv-123",
  "user_id": "user-entra-id",
  "tenant_id": "poc-tenant",
  "title": "TWI — CNC-01 gép beállítása",
  "content_type": "twi",
  "draft_content": "...",           // Utolsó jóváhagyott szöveg
  "pdf_blob_name": "twi/conv-123/abc.pdf",
```

```

"pdf_url": "https://...",
"llm_model": "mistral-large-latest",
"revision_count": 1,
"status": "approved" | "draft" | "rejected",
"created_at": ISODate,
"approved_at": ISODate,
"approved_by": "user-entra-id"
}
// Index: { tenant_id: 1, created_at: -1 }
// Index: { conversation_id: 1 }

// Collection: audit_log
{
  "_id": ObjectId,
  "conversation_id": "conv-123",
  "user_id": "user-entra-id",
  "tenant_id": "poc-tenant",
  "channel": "msteams",
  "event_type": "twi_generated" | "twi_approved" | "twi_rejected" | "twi_revised",
  "intent": "generate_twi",
  "llm_model": "mistral-large-latest",
  "llm_tokens_input": 1250,
  "llm_tokens_output": 2800,
  "revision_count": 1,
  "pdf_blob_name": "twi/conv-123/abc.pdf",
  "status": "completed",
  "approval_timestamp": ISODate,
  "created_at": ISODate
}
// Index: { tenant_id: 1, created_at: -1 }
// Index: { event_type: 1 }

```

7. PDF GENERÁLÁS

7.1 Pipeline

TWI szöveg (markdown-szerű)

- Jinja2 template rendering (HTML)
- WeasyPrint (HTML → PDF)
- Blob Storage upload
- URL visszaadás

7.2 Jinja2 HTML Template

html

```
<!-- app/templates/twi_template.html -->
<!DOCTYPE html>
<html lang="hu">
<head>
<meta charset="UTF-8">
<style>
@page {
    size: A4;
    margin: 2cm;
    @bottom-center {
        content: "agentize.eu — AI által generált tartalom — " counter(page) "/" counter(pages);
        font-size: 8pt;
        color: #888;
    }
}
body {
    font-family: Arial, sans-serif;
    font-size: 11pt;
    line-height: 1.5;
    color: #2c3e50;
}
.header {
    border-bottom: 3px solid #1b4f72;
    padding-bottom: 10px;
    margin-bottom: 20px;
}
.header h1 { color: #1b4f72; margin: 0; font-size: 18pt; }
.header .meta { color: #666; font-size: 9pt; margin-top: 5px; }
.ai-warning {
    background: #fef5e7;
    border-left: 4px solid #e67e22;
    padding: 10px 15px;
    margin: 15px 0;
    font-size: 9pt;
    color: #856404;
}
h2 { color: #2e86c1; font-size: 14pt; border-bottom: 1px solid #ddd; padding-bottom: 5px; }
h3 { color: #1b4f72; font-size: 12pt; }
.step {
    background: #f8f9fa;
    border: 1px solid #dee2e6;
    border-radius: 4px;
    padding: 12px;
    margin: 10px 0;
}
.step-number { color: #1b4f72; font-weight: bold; font-size: 13pt; }
```

```

.key-point { color: #2e86c1; font-style: italic; }

.reason { color: #666; font-size: 10pt; }

.approval-box {
    border: 2px solid #27ae60;
    background: #eafaf1;
    padding: 15px;
    margin-top: 30px;
}

.approval-box .label { font-weight: bold; color: #27ae60; }

</style>
</head>
<body>

<div class="header">
    <h1>{{ title }}</h1>
    <div class="meta">
        Generálva: {{ generated_at }} | Modell: {{ model }} | Verzió: {{ revision }}
    </div>
</div>

<div class="ai-warning">
     AI által generált tartalom — emberi felülvizsgálat szükséges.
    Ez a dokumentum az agentize.eu AI platform segítségével készült.
</div>

{{ content_html }}

{%- if approved %}

<div class="approval-box">
    <div class="label">  Jóváhagyva</div>
    <div>{{ approved_by }} — {{ approved_at }}</div>
</div>

{%- endif %}

</body>
</html>

```

7.3 PDF Generator Tool

python

```

# app/agent/tools/pdf_generator.py
import io
from datetime import datetime, timezone
from jinja2 import Environment, FileSystemLoader
from weasyprint import HTML
import markdown

template_env = Environment(loader=FileSystemLoader("app/templates"))

async def generate_twi_pdf(content: str, metadata: dict, user_id: str) -> bytes:
    """TWI tartalom → PDF bytes"""
    # Markdown → HTML konverzió
    content_html = markdown.markdown(content, extensions=["tables", "fenced_code"])

    # Jinja2 template rendering
    template = template_env.get_template("twi_template.html")
    html_content = template.render(
        title=_extract_title(content),
        generated_at=metadata.get("generated_at", "N/A"),
        model=metadata.get("model", "N/A"),
        revision=metadata.get("revision", 0),
        content_html=content_html,
        approved=True,
        approved_by=user_id,
        approved_at=datetime.now(timezone.utc).strftime("%Y-%m-%d %H:%M UTC"),
    )

    # HTML → PDF
    pdf_bytes = HTML(string=html_content).write_pdf()
    return pdf_bytes

def _extract_title(content: str) -> str:
    """Első sor kinyerés címként"""
    for line in content.split("\n"):
        line = line.strip().lstrip("#").strip()
        if line and not line.startswith("⚠️"):
            return line[:100]
    return "TWI Munkautasítás"

```

8. AZURE SZOLGÁLTATÁS KLIENSEK

8.1 AI Foundry Client

python

```
# app/services/ai_foundry.py
import logging
from azure.ai.inference import ChatCompletionsClient
from azure.core.credentials import AzureKeyCredential

from app.config import settings

logger = logging.getLogger(__name__)

_client = None

def _get_client():
    global _client
    if _client is None:
        _client = ChatCompletionsClient(
            endpoint=settings.ai_foundry_endpoint,
            credential=AzureKeyCredential(settings.ai_foundry_key),
        )
    return _client

async def call_llm(
    prompt: str,
    system_prompt: str = None,
    temperature: float = None,
    max_tokens: int = None,
) -> str:
    """LLM hívás az Azure AI Foundry-n keresztül"""
    client = _get_client()

    messages = []
    if system_prompt:
        messages.append({"role": "system", "content": system_prompt})
    messages.append({"role": "user", "content": prompt})

    response = client.complete(
        messages=messages,
        model=settings.ai_model,
        temperature=temperature or settings.ai_temperature,
        max_tokens=max_tokens or settings.ai_max_tokens,
    )

    result = response.choices[0].message.content

    # Token tracking (audit-hoz)
    usage = response.usage
```

```
logger.info(  
    f"LLM call: model={settings.ai_model}, "  
    f"input_tokens={usage.prompt_tokens}, "  
    f"output_tokens={usage.completion_tokens}"  
)  
  
return result
```

8.2 Cosmos DB Client

```
python
```

```
# app/services/cosmos_db.py
from motor.motor_asyncio import AsyncIOMotorClient
from datetime import datetime, timezone

from app.config import settings

_client = None
_db = None

def _get_db():
    global _client, _db
    if _db is None:
        _client = AsyncIOMotorClient(settings.cosmos_connection)
        _db = _client[settings.cosmos_database]
    return _db

class ConversationStore:
    def __init__(self):
        self.collection = _get_db()["conversations"]

    async def get_or_create(self, conversation_id: str, user_id: str, channel: str, tenant_id: str = "poc-tenant"):
        doc = await self.collection.find_one({"conversation_id": conversation_id})
        if doc:
            await self.collection.update_one(
                {"conversation_id": conversation_id},
                {"$set": {"last_activity": datetime.now(timezone.utc)}, "$inc": {"message_count": 1}}
            )
        return doc

        new_doc = {
            "conversation_id": conversation_id,
            "user_id": user_id,
            "tenant_id": tenant_id,
            "channel": channel,
            "started_at": datetime.now(timezone.utc),
            "last_activity": datetime.now(timezone.utc),
            "message_count": 1,
            "status": "active",
        }
        await self.collection.insert_one(new_doc)
        return new_doc

class AuditStore:
    def __init__(self):
```

```
self.collection = _get_db()["audit_log"]

async def log(self, entry: dict):
    entry["created_at"] = datetime.now(timezone.utc)
    await self.collection.insert_one(entry)

class DocumentStore:
    def __init__(self):
        self.collection = _get_db()["generated_documents"]

    async def save(self, doc: dict):
        doc["created_at"] = datetime.now(timezone.utc)
        await self.collection.insert_one(doc)
        return doc
```

8.3 Blob Storage Client

python

```

# app/services/blob_storage.py
from azure.storage.blob import BlobServiceClient, generate_blob_sas, BlobSasPermissions
from datetime import datetime, timezone, timedelta

from app.config import settings

_client = None

def _get_client():
    global _client
    if _client is None:
        _client = BlobServiceClient.from_connection_string(settings.blob_connection)
    return _client

async def upload_pdf(pdf_bytes: bytes, blob_name: str) -> str:
    """PDF feltöltés Blob Storage-ba, SAS URL visszaadás"""
    client = _get_client()
    container_client = client.get_container_client(settings.blob_container)
    blob_client = container_client.get_blob_client(blob_name)

    blob_client.upload_blob(pdf_bytes, content_settings={"content_type": "application/pdf"})

    # SAS token generálás (24 órás érvényesség)
    sas_token = generate_blob_sas(
        account_name=client.account_name,
        container_name=settings.blob_container,
        blob_name=blob_name,
        account_key=client.credential.account_key,
        permission=BlobSasPermissions(read=True),
        expiry=datetime.now(timezone.utc) + timedelta(hours=24),
    )

    return f'{blob_client.url}?{sas_token}'

```

9. DOCKER + CI/CD

9.1 Dockerfile

dockerfile

```
FROM python:3.12-slim

# WeasyPrint rendszer-függőségek
RUN apt-get update && apt-get install -y --no-install-recommends \
    libpango-1.0-0 \
    libpangoft2-1.0-0 \
    libharfbuzz0b \
    libffi-dev \
    libgdk-pixbuf2.0-0 \
    && rm -rf /var/lib/apt/lists/*

WORKDIR /app

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY app/ ./app/

EXPOSE 8000

CMD ["uvicorn", "app.main:app", "--host", "0.0.0.0", "--port", "8000", "--workers", "2"]
```

9.2 Dev Container

json

```
// .devcontainer/devcontainer.json
{
  "name": "agentize-poc",
  "image": "mcr.microsoft.com/devcontainers/python:3.12",
  "features": {
    "ghcr.io/devcontainers/features/azure-cli:1": {},
    "ghcr.io/devcontainers/features/docker-in-docker:2": {}
  },
  "postCreateCommand": "pip install -r requirements.txt",
  "customizations": {
    "vscode": {
      "extensions": [
        "ms-python.python",
        "ms-azurertools.vscode-bicep",
        "ms-azurertools.vscode-azureresourcegroups"
      ]
    }
  },
  "forwardPorts": [8000]
}
```

9.3 GitHub Actions (Minimal PoC)

yaml

```
# .github/workflows/deploy.yml
name: Deploy PoC Backend

on:
  push:
    branches: [main]

jobs:
  build-and-deploy:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4

      - name: Login to Azure
        uses: azure/login@v2
        with:
          creds: ${{ secrets.AZURE_CREDENTIALS }}

      - name: Build and push to ACR
        run:
          az acr build \
            --registry ${{ vars.ACR_NAME }} \
            --image poc-backend:${{ github.sha }} \
            --file Dockerfile .

      - name: Deploy to Container App
        run:
          az containerapp update \
            --name ca-agentize-poc-backend \
            --resource-group rg-agentize-poc-swedencentral \
            --image ${{ vars.ACR_NAME }}.azurecr.io/poc-backend:${{ github.sha }}
```

10. TEAMS APP MANIFEST

json

```
// teams-app/manifest.json
{
  "$schema": "https://developer.microsoft.com/en-us/json-schemas/teams/v1.17/MicrosoftTeams.schema.json",
  "manifestVersion": "1.17",
  "version": "0.1.0",
  "id": "{{BOT_APP_ID}}",
  "developer": {
    "name": "agentize.eu",
    "websiteUrl": "https://agentize.eu",
    "privacyUrl": "https://agentize.eu/privacy",
    "termsOfUseUrl": "https://agentize.eu/terms"
  },
  "name": {
    "short": "agentize AI",
    "full": "agentize.eu Enterprise AI Platform"
  },
  "description": {
    "short": "AI munkautasítás generátor",
    "full": "Enterprise AI platform gyártási munkautasítások (TWI) generálásához. GDPR-konform, EU adatközpontban fut."
  },
  "icons": {
    "outline": "outline.png",
    "color": "color.png"
  },
  "accentColor": "#1B4F72",
  "bots": [
    {
      "botId": "{{BOT_APP_ID}}",
      "scopes": ["personal", "team", "groupChat"],
      "supportsFiles": false,
      "commandLists": [
        {
          "scopes": ["personal"],
          "commands": [
            {
              "title": "Új TWI",
              "description": "Új munkautasítás generálása"
            },
            {
              "title": "Segítség",
              "description": "Használati útmutató"
            }
          ]
        }
      ]
    }
  ]
}
```

```
],  
  "permissions": ["identity", "messageTeamMembers"],  
  "validDomains": ["{{BACKEND_FQDN}}"]  
}
```

11. KÖRNYEZETI VÁLTOZÓK

```
bash  
  
# .env.example  
  
# Azure AI Foundry  

```

12. IMPLEMENTÁCIÓS SORREND — 10 MUNKANAP

Nap 1-2: Infrastruktúra

- Bicep template megírás (main.bicep + parameters.json)
- `(az deployment group create)` — test deploy Sweden Central-ba
- AI Foundry model deployment (Mistral Large, Data Zone Standard)
- Entra ID App Registration (Bot Framework-höz)
- Key Vault secrets feltöltés
- Cosmos DB collections + indexek létrehozás
- Blob Storage container létrehozás
- Validáció:** minden resource elérhető, AI Foundry válaszol

Nap 3-4: Backend Core

- Projekt struktúra létrehozás (lásd 4.1)
- `(app/config.py)` — Settings osztály
- `(app/services/ai_foundry.py)` — LLM client + test hívás
- `(app/services/cosmos_db.py)` — MongoDB client + test CRUD
- `(app/agent/state.py)` — AgentState definíció
- `(app/agent/nodes/intent.py)` — Intent recognition
- `(app/agent/nodes/generate.py)` — TWI generálás
- `(app/agent/graph.py)` — Alap gráf (intent → generate → END)
- Validáció:** `(pytest)` — intent felismerés + generálás működik

Nap 5-6: Bot Framework + Adaptive Cards

- `(app/main.py)` — FastAPI + /api/messages endpoint
- `(app/bot/bot_handler.py)` — Message handler
- `(app/bot/adaptive_cards.py)` — Review + Approval + Result + Welcome cards
- Azure Bot Service resource regisztráció
- Teams channel + Telegram channel bekapcsolás
- Teams App Manifest (sideload-hoz)
- Validáció:** Teams-ből üzenet → Bot válaszol Adaptive Card-dal

Nap 7-8: Human-in-the-loop + Revision Loop

- `(app/agent/nodes/review.py)` — Review checkpoint
- `(app/agent/nodes/revise.py)` — Revision node
- `(app/agent/nodes/approve.py)` — Final approval checkpoint
- Graph interrupt_before konfigurálás
- Bot handler: Adaptive Card submit → graph resume
- Revision loop tesztelés (max 3 kör)
- Validáció:** Teljes flow: generálás → review → szerkesztés → jóváhagyás

Nap 9: PDF + Output

- `[app/templates/twi_template.html]` — Jinja2 template
- `[app/agent/tools/pdf_generator.py]` — PDF generálás (WeasyPrint)
- `[app/services/blob_storage.py]` — PDF upload + SAS URL
- `[app/agent/nodes/output.py]` — Output node
- `[app/agent/nodes/audit.py]` — Audit log node
- Dockerfile + local build + test
- Validáció:** PDF generálódik, letölthető URL-ről, audit log Cosmos DB-ben

Nap 10: Integration Test + Demo Prep

- End-to-end teszt: Teams → Bot → Agent → PDF → letöltés
 - End-to-end teszt: Telegram → Bot → Agent → PDF → letöltés
 - Container App deployment (GitHub Actions vagy manuális)
 - Demo script megírás (3 perces happy path)
 - Known issues dokumentálás
 - Validáció:** Demo futtatás valós Teams + Telegram környezetben
-

13. DEMÓ SCRIPT (3 perc)

1. [Teams-ben] "Szia! Készíts egy TWI utasítást a CNC-01 gép napi karbantartásáról."
→ Bot: "🕒 Feldolgozom..."
→ Bot: [Review Adaptive Card — vázlat megjelenik]
2. [Review Card-on] "📝 Szerkesztés kérem" + "A 3. lépéskor add hozzá a hőmérséklet ellenőrzést!"
→ Bot: "🕒 Módosítom..."
→ Bot: [Módosított Review Adaptive Card]
3. [Review Card-on] "✓ Jóváhagyom a vázlatot"
→ Bot: [Final Approval Card]
4. [Approval Card-on] "✓ Ellenőriztem és jóváhagyom"
→ Bot: "🕒 PDF generálás..."
→ Bot: [Result Card — PDF letöltés link]
5. PDF megnyitás — formázott dokumentum, agentize.eu branding, AI jelölés, jóváhagyási info
6. [Opcionális — Telegram-ban ugyanez]

14. SIKERESSÉGI KRITÉRIUMOK

A PoC akkor kész, ha:

1. Bicep template egy gombnyomásra települ Sweden Central-ba
 2. Teams-ből végigfut a teljes flow: kérés → generálás → review → szerkesztés → jóváhagyás → PDF
 3. Telegram-ból ugyanez működik
 4. PDF letölthető, formázott, tartalmazza az AI jelölést és jóváhagyási infót
 5. Audit log Cosmos DB-ben: ki, mit, mikor, melyik modellel
 6. EU AI Act jelölés minden AI output-on
 7. Többpontos jóváhagyás működik (review + final approval)
 8. Revision loop működik (szerkesztés kérés → módosítás → újra review)
 9. 3 perces demó végigfuttatható hiba nélkül
-